

Implementing Digital Twins with Elixir and OTP

Executive Summary

Digital twins are live virtual replicas of physical systems or processes, continuously synchronized with real-world data. They require high concurrency, real-time updates, and robust fault tolerance. Elixir (on the Erlang VM) and its OTP framework offer built-in support for massive concurrency, fault-isolated processes, distribution, and live code upgrades, making them a strong fit for digital twin platforms. This white paper analyzes how digital twin requirements map to OTP primitives (GenServer, Supervisors, Registry, DynamicSupervisor, GenStateMachine, etc.) and OTP services (clustering, Mnesia, ETS, Telemetry, releases). We compare Elixir/OTP to alternatives (Java/Akka, Go, Node.js, Rust) in terms of latency, throughput, state management, and developer productivity. Case studies—from industrial IoT to healthcare—demonstrate Erlang/Elixir in similar domains. We provide architecture and sequence diagrams, recommended patterns, code snippets, testing/benchmarking strategies, and discuss deployment, observability, security, and cost/team considerations. In summary, leveraging Elixir/OTP can yield highly concurrent, resilient digital twins with hot-upgrade capabilities and efficient use of resources [26†L19-L23] [15†L128-L137] .

What Is a Digital Twin and Key Requirements

A **digital twin** is a virtual model of a physical object or system that is kept in sync with the real-world counterpart via live data feeds. McKinsey defines it as “a digital replica of a physical object, person, or process ... linked to real data sources ... updating in real time to reflect the original” [5†L41-L46] . Digital twins enable simulation, analysis, and optimization (e.g. “what-if” scenarios) without risk to the physical asset.

Key requirements for digital twins include:

- **Real-Time Data Integration:** Continuous ingestion of sensor or event data (IoT telemetry, control signals) to update the twin.
- **High Concurrency:** Each twin may be modeled as an independent process; systems often manage many twins simultaneously.
- **Isolation & Fault-Tolerance:** Failures in one twin should not crash the system. Supervisors must auto-restart failed components.
- **State Management:** Twins often maintain rich state (historical logs, simulation state). This may require in-memory state, databases, or hybrid approaches.
- **Scalability:** The system must scale to potentially thousands or millions of twins (e.g. IoT devices, production lines). Clustering across nodes is essential for large-scale deployments.
- **Observability:** Monitoring and metrics (latency, throughput, resource usage) are crucial to validate twin fidelity and performance.
- **Hot Code Updates:** Digital twins may require live updates (e.g. changing models) without downtime. OTP's hot code swapping can be leveraged.

- **Security and Consistency:** Secure communication and consistent state (especially under partitions) are critical, sometimes implying transactional updates (e.g. using Erlang Mnesia with majority commit).

These requirements align closely with the capabilities of the Erlang/Elixir runtime: it excels at concurrent, distributed, fault-tolerant applications [15†L128-L137] [26†L19-L23] . In Elixir/OTP, each twin can be implemented as an isolated process (GenServer or GenStateMachine), supervised for restart logic, and clustered transparently across nodes. Erlang's lightweight processes and preemptive scheduling support millions of concurrent twins if needed [26†L19-L23] [15†L128-L137] .

Mapping Digital Twin Architecture to OTP Primitives


Elixir/OTP provides a rich toolbox for building the components of a digital twin system. Table 1 maps common digital twin components to OTP abstractions:

- **Twin Process (stateful actor):** Each digital twin instance can be a `GenServer` or `GenStateMachine` . The `GenServer` behavior encapsulates a stateful process with `handle_call` / `handle_cast` callbacks [18†L18-L24] . For modeling state machines or workflows, `GenStateMachine` allows defining state-transition logic with callbacks [41†L12-L16] .
- **Supervision Tree:** A Supervisor process supervises child processes. Supervisors create a **supervision tree** providing fault tolerance [31†L18-L24] . For example, a root supervisor may spawn a `DynamicSupervisor` for twins and other supervisors for ingestion, web interface, etc. Supervisors use restart strategies (`:one_for_one` , `:rest_for_one` , etc.) to control how children are restarted on failure [65†L44-L52] [31†L18-L24] .
- **Dynamic Supervisor:** Use `DynamicSupervisor` to manage a dynamic set of twin processes. A `DynamicSupervisor` starts with no children and spawns them on demand via `DynamicSupervisor.start_child` [35†L14-L21] . This is ideal when twin processes are created/terminated at runtime (e.g. new physical assets coming online).
- **Registry (Process Lookup):** The `Registry` module is a decentralized key-value process store [33†L18-L23] . Each twin process can register itself with a unique key (e.g. asset ID), making it discoverable by other processes or APIs via `Registry.lookup` [33†L18-L27] . If a twin process crashes, its registry entry is removed automatically [33†L25-L28] , avoiding stale references.
- **OTRelease/Deployment:** Elixir releases (e.g. via `mix release`) bundle the BEAM VM and application code for deployment. OTP releases support versioned upgrades, allowing hot code swapping when configured. This supports rolling updates of twin logic without downtime.
- **Clustering & Distribution:** Erlang's VM supports clustering out-of-the-box [29†L47-L55] . Multiple BEAM nodes can form a cluster, and processes can message each other transparently across nodes. This enables distributing twins across machines. OTP also supports node discovery (via `libcluster` or custom) for dynamic clusters.
- **State Storage (Persistence):** For persistence and shared state, Erlang provides **ETS** (in-memory tables) and **Mnesia** (distributed DB). Mnesia allows tables to be replicated on RAM or disk across nodes, supporting transactions and majority commit [43†L22-L30] . For critical twin state, Mnesia can ensure consistency and failover. ETS can serve as a fast cache or session store if persistence across restarts is not needed.
- **Telemetry/Monitoring:** Elixir's `:telemetry` library lets components emit metrics/events (e.g. message latencies, process counts). These can be aggregated and exported to monitoring systems

(Prometheus, Grafana). OTP's built-in tools like Observer (for live introspection) and Logger (for logs) add observability.

- **UI/Dashboard (LiveView):** If the digital twin has a web dashboard, Phoenix LiveView can push real-time updates to browser clients. LiveView integrates with Telemetry and PubSub for pushing twin state changes to a UI with minimal JS.
- **Connectivity (IoT protocols):** Erlang/Elixir has libraries for MQTT (e.g. EMQX), Zigbee, BLE, etc. thanks to BEAM's robust binary handling [15†L110-L119] . A digital twin platform often sits behind such message brokers. Elixir's ease of handling binary protocols simplifies integration with IoT devices.

By leveraging these primitives, an Elixir/OTP system for digital twins might have the following outline:



This sequence shows how sensor data flows through a broker into an ingestion process, triggers dynamic creation of a twin actor (supervised and registered), which updates its state and notifies observers.

Concurrency and Fault Tolerance

Erlang/OTP's concurrency model is built on lightweight isolated processes (actors). Each process has its own memory and mailbox, eliminating shared-memory conflicts. Processes are extremely cheap: benchmarks show that Erlang can spawn millions of processes; in one study 9 million Erlang processes were spawned in ~58 seconds [26†L19-L23] . For latency-critical messaging, both Erlang and Go perform best, outpacing frameworks like Akka [26†L19-L23] . In practice, OTP systems routinely handle hundreds of thousands of actors (e.g. WhatsApp's Erlang servers handle billions of messages/day [26†L61-L68]).

Fault Isolation: Since each twin runs in its own process, a crash affects only that twin. Supervisors restart failed child processes according to policy [31†L18-L24] . For example, a `:one_for_one` supervisor restarts a crashed twin without affecting siblings [65†L44-L52] . Supervisors themselves can be supervised (hierarchies of supervisors), yielding "let it crash" architectures where faulty components are healed automatically. The official docs note: "Supervision trees provide fault-tolerance and encapsulate how our applications start and shutdown" [31†L18-L24] . If a twin's GenServer handler crashes on unexpected input, its supervisor will restart it cleanly (as shown with the Counter example resetting to 0) [31†L113-L117] .

Scalability and Distribution: OTP's built-in distribution means that processes can communicate across nodes with the same APIs as local calls. A clustered Elixir system can place twin processes on multiple machines to balance load. The BEAM VM handles distribution: "If you are newer to Elixir, clustering isn't creating a web API – it's connecting multiple Elixir/Erlang runtime systems over a network into a single distributed system... a core part of the language and OTP framework" [29†L47-L55] . This makes it straightforward to scale horizontally. Libraries like libcluster can automate node discovery in cloud environments. Telemetry and distributed PubSub (Phoenix.PubSub or Registry.dispatch) allow state or events to be shared cluster-wide.

Data Consistency: For state synchronization, Mnesia's distributed tables support transactions and majority safety [43†L22-L30] . You can designate mission-critical tables with `majority = true` so that a majority of replicas must agree before commit, avoiding split-brain inconsistencies. Alternatively, Event Sourcing patterns (e.g. persisting events in ETS/Mnesia or Kafka) ensure that a twin's state can be rebuilt after failure. ETS can act as a fast in-memory cache for recent state or indexes, while the durable state resides in Mnesia or external DB. In practice, OTP teams often use the BEAM's strengths – fast messaging and cheap processes – for stateful event-driven models rather than heavy DB access for every update [72†L276-L284] .

Hot Code Upgrade and Releases

An important operational feature of OTP is *hot code swapping*. OTP releases (via `mix release` or Distillery) package the application and the BEAM into a self-contained artifact. Releases support upgrades: new code versions can be loaded into a running system via `:sys.change_code` or `code_change` callbacks. Although implementing code upgrades requires planning (modules can implement a `code_change/4` function to migrate state), it allows zero-downtime updates of twin logic. For example, if twin algorithms change, the old GenServer versions can transform their state and continue under new code.

A release also simplifies deployment: it embeds the Erlang runtime so the system can run independently of external dependencies [29†L47-L55] . Tools like Mix Releases and Distillery also allow custom boot scripts and hooks. Combined with containerization (Docker) or system packages, releases make it easy to distribute and launch Elixir apps.

In short, OTP's release handling means updates and rollbacks of the twin platform can be done live, preserving existing processes and state (subject to the `code_change` implementation). This advantage is rare in other ecosystems.

Observability and Telemetry

Elixir/OTP provides rich observability tools. Developers can instrument processes and applications with the [Telemetry library](#) (emit metrics on events, function calls, process lifecycle, etc.). Telemetry metrics (counters, histograms) can be consumed by reporters for Prometheus/Grafana. For example, each twin GenServer can emit events on state updates or errors, which can be aggregated and visualized.

The Observer tool (built into Erlang) offers a live GUI to inspect a running system: view process lists, memory, supervision trees, message queue lengths, and more. This can be invaluable in development and debugging. Logging via Elixir's `Logger` and tracing (via `:dbg` or `:observer_cli`) adds depth. Telemetry is

already integrated into Phoenix (1.5+) and many libraries, making it easy to treat metrics as a first-class citizen **【19†L5-L9】** .

Live dashboards can be built with Phoenix/LiveView to display twin metrics in real time. For example, a LiveView page can subscribe to a PubSub topic that receives Telemetry events from all twin processes (or periodic snapshots of their state), updating graphs or maps instantly in the browser without custom JavaScript.

Alternative Platforms Comparison

Feature	Elixir/Erlang (BEAM) 【26†L19-L23】 【15†L128-L137】	Java/Scala + Akka	Go	Node.js	Rust
Concurrency Model	Actor-based (processes), lightweight (~0.5μs spawn) 【26†L19-L23】	Actors (Akka) on JVM threads	Goroutines (M:N multiplexing)	Event loop + async	Threads/futures, no GC overhead
Latency (msg)	Very low latency (Erlang ranks best for msg latency) 【26†L19-L23】	Akka medium (Scala), but lower with Akka	Low latency (Go) 【26†L19-L23】	Higher (single thread, callback)	Low (compiled, no GC during computation)
Throughput	High (comparable to Scala/Akka); lower than Go for raw compute 【26†L19-L23】	High; Akka can handle many actors	Very high (compiled, optimized) 【26†L19-L23】	Moderate (single core, single thread)	Very high (native speed, safe concurrency)

Feature	Elixir/Erlang (BEAM) 【26†L19-L23】 【15†L128-L137】	Java/Scala + Akka	Go	Node.js	Rust
State Management	Built-in per-process state + Mnesia/ETS (distributed DBs) 【43†L22-L30】	Akka has Persistence (via event logs or DBs)	Sync/async with external DB (no built-in actor)	Typically use external DB (Redis, etc)	no built-in; use crates (sled, or external DB)
Fault Isolation	Strong: process crash isol., supervisors restart 【31†L18-L24】	Good: actor restarts, but on JVM one actor crash can affect thread pool	Partial: panic in goroutine kills only that goroutine (others run)	Weak: uncaught error can crash process	Good: thread panic unwinds that thread, others safe (plus <code>catch_unwind</code>)
Hot Upgrades	Supported (OTP releases, <code>code_change</code>)	Not natively (requires complex setup)	Not supported (binary reload not available)	Not supported (need restart)	Not supported (rust binaries static)
Dev Productivity	High-level functional, dynamic typing; rapid iteration	Statically typed; verbose in Akka/Java, moderate in Scala	Simple syntax, but manual concurrency	JS very quick to start, but concurrency via callbacks/promises	Steep learning curve, strict; memory safety yields reliability
Ecosystem Maturity	Mature for telecom/messaging/IoT; growing web/ecosystem	Very mature; JVM ecosystem, rich libraries	Growing; strong in web/tools/microservices	Very mature (npm ecosystem) but less concurrency focus	Emerging (crates.io growing); less libraries for higher-level services
Operational Tooling	Built-in clustering, releases, Observer; dist tooling	Java tooling (JVM monitoring, Docker) plus Akka Clustering	Good; container-friendly, many ORMs, metrics libs	Mature DevOps ecosystem, Node monitors	Container-friendly, cargo build, less runtime support

Feature	Elixir/Erlang (BEAM) 【26†L19-L23】 【15†L128-L137】	Java/Scala + Akka	Go	Node.js	Rust
Team Skill / Cost	Niche language: smaller talent pool; training needed; but high reliability reduces maintenance. Low CPU/RAM usage can reduce infra cost.	Large talent pool; JVM runtime overhead; enterprise adoption	Popular; moderate talent; compiled overhead moderate; static compilation ease ops.	Huge talent pool; single-thread limitations; high-per-node usage possible.	Growing talent; high dev cost (strict), but low-run cost (fast).

Sources: Valkov et al. (2018) compare Erlang, Go, Scala/Akka on concurrency benchmarks 【26†L19-L23】 . Erlang/Elixir shines in process spawn latency and handling short-lived actors, while Go wins throughput. Akka (Scala) can host the most actors but with higher latency. These findings inform our table: Elixir (Erlang) and Go excel at low-latency messaging, Go leads raw throughput, and Akka excels in large process counts 【26†L19-L23】 . Node.js, being single-threaded, generally lags in raw concurrency for CPU-bound tasks, though it's effective for I/O-bound workloads. Rust offers low-level control and very high throughput but lacks the dynamic actor ecosystem out-of-the-box (though actor crates are emerging). Elixir/OTP's mature actor tools, built-in distribution, and fault model give it an edge for complex, stateful twin systems, at the cost of requiring BEAM expertise.

Case Studies and Examples

- **Industrial IoT & Automation:** Erlang Solutions notes that the BEAM is ideal for IoT gateways and brokers. It “provides all the tools to easily tackle new and existing protocols” (handling binary IoT protocols like Zigbee, BLE, etc.) and its “low latency, fault tolerance, and massive concurrency make it well-suited for deploying IoT solutions at an industrial scale” 【15†L110-L119】 【15†L128-L137】 . In practice, Erlang-based brokers (e.g. EMQX) are used in large IoT deployments.
- **Telecom / Messaging:** WhatsApp famously uses Erlang for its chat servers (handling billions of messages/day) 【26†L61-L68】 . This shows BEAM's ability to maintain stateful connections at scale. While not a “twin” of a physical asset, the use-case is analogous: each user/session is an actor process with high availability.
- **Automotive Digital Twins:** RemotiveLabs built an automotive development platform (virtual ECUs & signals) on Elixir. They report that “Elixir already has a toolbox... to handle many signals in parallel – also in automotive” and OTP provides a “stable architecture” with isolated processes that restart independently 【46†L81-L90】 . This platform (RemotiveBroker) models virtual vehicle components and supports continuous integration for vehicle software, effectively acting as a system-level twin of the vehicle electronics.

- **Healthcare (COVID Simulation):** The HCA Waterpark project modeled each patient as a GenServer (“patient actor”) running from admission to discharge, a form of digital twin of the patient’s state [72†L276-L284] . Each patient process accumulates all HL7 messages for that patient (simulating history), and GenServers hibernate when idle and use eviction policies to conserve memory [72†L288-L292] . This demonstrates long-lived stateful actors with fault tolerance (continuity of care across system restarts) and efficient state management – precisely what OTP provides.

【75†embed_image】 *Figure: Slide from HCA’s Waterpark (COVID) project: each patient is managed by an Erlang process (“digital twin”), not just a database row* [72†L276-L284] .

These cases illustrate that Erlang/Elixir is used where real-time reliability and high concurrency matter. The patient-actor and RemoteLabs examples show digital-twin-like architectures: dynamic processes mirroring real-world entities, handled with OTP supervision and messaging.

Architecture Diagrams (Mermaid)

Below is a sample high-level component diagram for a digital twin system:

```

flowchart LR
    subgraph DataSources [Physical/IoT Layer]
        Sensor1["Sensor 1"]
        Sensor2["Sensor 2"]
    end
    subgraph Edge ["Edge/Cloud Services"]
        Broker[Message Broker]
        Ingest[Ingest Service (GenServer)]
        TwinManager[DynamicSupervisor<br/>Twins Manager]
        Registry[Registry]
        TwinA[Twin Process A]
        TwinB[Twin Process B]
        UI[Phoenix LiveView UI]
    end
    Sensor1 -->|data| Broker --> Ingest
    Sensor2 -->|data| Broker --> Ingest
    Ingest -->|start twin| TwinManager
    TwinManager --> TwinA & TwinB
    TwinA -->|register| Registry
    TwinB -->|register| Registry
    UI -->|user query| Registry --> TwinA
    Registry -->|route| TwinA -->|state| UI

```

This diagram shows IoT sensors feeding data to an ingestion GenServer, which under a DynamicSupervisor spawns twin processes. Each twin registers itself in a Registry for lookup. The UI (e.g. a LiveView dashboard) queries the Registry to interact with twin processes.

For sequence flows, the earlier mermaid diagram captures the lifecycle of data and twin updates (in text form above). These mermaid diagrams can be rendered into static images if needed; for brevity we use text.

Implementation Patterns & Code Snippets

- **Process-per-Twin:** Define a `GenServer` or `GenStateMachine` module (e.g. `Twin`) to hold one twin's state. The `init/1` callback sets initial state. Example initialization:

```
defmodule Twin do
  use GenServer
  def start_link(id), do: GenServer.start_link(__MODULE__, id, name:
via_tuple(id))
  def init(id), do: {:ok, %{id: id, data: nil}}
  # handle_call / handle_cast for operations...
end
```

Use `{:via, Registry, {MyRegistry, id}}` as the `name` if you need Registry lookup [33†L18-L23] .

- **Supervisor Setup:** In your Application's `start/2`, set up a tree:

```
children = [
  {Registry, keys: :unique, name: MyApp.Registry},
  {DynamicSupervisor, strategy: :one_for_one, name: MyApp.TwinSupervisor}
]
Supervisor.start_link(children, strategy: :one_for_all)
```

Then start twins dynamically:

```
DynamicSupervisor.start_child(MyApp.TwinSupervisor, {Twin, twin_id})
```

- **State Machine Example:** If the twin has modes (e.g. Active/Inactive), use `GenStateMachine` :

```
defmodule DeviceTwin do
  use GenStateMachine, callback_mode: :state_functions
  # state functions: off/3 and on/3
  def off(:cast, :activate, data) do {:next_state, :on, data} end
  def on(:cast, :deactivate, data) do {:next_state, :off, data} end
  # ... handle other messages
end
```

- **Supervision Strategies:** Use `:one_for_one` for independent twins; `:rest_for_one` or `:one_for_all` if twins share resources that require coordinated restarts [65†L44-L52] .

- **Telemetry Instrumentation:** Emit Telemetry events in the GenServer callbacks (e.g. on state update or receive). Then attach reporters:

```
:telemetry.execute([:twin, :update], %{count: 1}, %{twin_id: id})
```

- **Hot-Upgrade (Advanced):** Implement a `code_change/4` callback in your GenServer/GenStateMachine to transform state when you upgrade. Alternatively use Erlang's `hot_upgrade` tools with OTP releases.
- **Testing:** Use `ExUnit` for unit tests of GenServers (you can start them with `start_supervised`). For performance, use `Benchee` or `BencheeGraph` to benchmark message passing or twin throughput. Simulate load by spawning many GenServers or sending many messages.

Testing and Benchmarking

For correctness, write `ExUnit` tests for each GenServer callback, simulating messages (calls/casts) and verifying state transitions. Use property-based testing (PropCheck or StreamData) to cover edge cases (e.g. message floods, error conditions). Test supervision by deliberately causing child failures and asserting restart.

For performance, benchmark critical paths. For example, use [Benchee](#) to measure message latency and throughput: spawn many GenServers, send asynchronous casts (`GenServer.cast/2`) at varying rates, measure latency. Compare with alternatives where needed. Valkov et al.'s benchmarking approach [26†L19-L23] can guide: create microbenchmarks for spawn time, ping-pong messaging, etc. Also instrument in-production telemetry (via Prometheus) to catch bottlenecks (e.g. if a twin takes too long to process a message).

Security, Data Consistency, and Persistence

- **Security:** Use TLS for any external connections (e.g. TLS with MQTT). Within a cluster, use Erlang's distribution with short cookie or handshake. Run BEAM nodes in isolated environments (containers or VMs). Code correctness: Elixir's functional model and immutable data reduce certain vulnerabilities. For embedded or secure scenarios, consider running on secure BEAM variants (AtomVM for microcontrollers, or OTP on seL4). Always validate external data in GenServer handlers to avoid crashes or injection.
- **Data Consistency:** Mnesia transactions provide atomic multi-table updates [43†L22-L30]. For example, to update a twin's state and a summary table, wrap in `mnesia.transaction/1`. This avoids partial state. Be wary of network partitions: design your cluster (number of Mnesia copies, `majority` flag) to tolerate node loss without split-brain. Alternatively, use event logs (Kafka, RabbitMQ) plus eventual consistency if strict transactions are not needed.
- **Persistence:** Decide what state to persist in Mnesia vs. what to keep only in memory. For example, current twin status can be in-memory GenServer data, but periodic checkpoints or critical events can be recorded to Mnesia or disk. You can configure Mnesia tables with `disc_copies` for durability.

ETS can hold frequently accessed state in RAM, with Mnesia as persistent backing (ETS views of Mnesia tables are possible).

Deployment and Operations

Deploy Elixir/OTP releases on servers or containers. A common pattern is to use Docker images containing the release and then use orchestration (Kubernetes, ECS) for scaling and failover. Each BEAM node should be given a unique name (often via hostname or env var) and cookie for clustering. Configuration (node list, Mnesia schema) can be passed via environment or a KV store.

Elixir's OTP also integrates with container monitors. For zero-downtime, Kubernetes can do rolling upgrades of pods (each with a new release version). Intra-cluster sharding can distribute twins. Tools like `libcluster` can auto-connect nodes in the same network. Telemetry data (e.g. via `TelemetryMetrics`) can be exported to Prometheus.

Monitor VM-level stats via Observer or `:observer_cli` in dev, and in prod watch BEAM's memory and process count. The BEAM's memory limits can be tuned (`+S` flags, etc.) to prevent system exhaustion.

Migration Considerations

If migrating an existing twin system to Elixir/OTP, consider the following:

- **Data Model:** If coming from relational DB rows, you'll re-architect into actors (as HCA did). Each entity becomes a process; existing data becomes initial state. Migrate data by initializing `GenServers` from DB on startup or on first access.
- **Protocol Translation:** Map existing messaging interfaces (REST calls, MQTT, OPC-UA) to Elixir. Plenty of libraries exist, but custom adapters may be needed.
- **Training:** Elixir's concurrency model is different; teams must learn OTP patterns (supervision, message passing). However, this investment pays off in reliability. Elixir's syntax (similar to Ruby) and powerful tooling ease adoption.
- **Incremental Integration:** You can integrate Elixir components gradually. For example, start by using Phoenix channels or ports to communicate with legacy parts, then migrate critical functions to Elixir.
- **Ecosystem Gaps:** Some specialized domain libraries may not exist in Elixir. In such cases, interop with Erlang or NIFs is possible, but keep it minimal. For example, if heavy numerical simulation is needed, it might be done in C/Rust as a native port, with Elixir coordinating.

Cost and Team Impact

Budget and team size aside, consider:

- **Infrastructure:** Elixir's efficient use of hardware (Erlang processes use little memory, minimal GC pause) can reduce server costs. High availability features are built-in, potentially lowering the need for extra HA tools.

- **Developer Costs:** Elixir devs are less common than Java/JavaScript devs, which might increase hiring costs. However, Elixir's developer productivity (clear abstractions, fewer bugs in concurrency logic) can reduce time-to-market.
- **Maintenance:** Systems written in Elixir/OTP can run for years without restart, reducing operational churn. Fault isolation means single errors are less likely to require human intervention.
- **Training:** Existing teams will need training in Elixir and functional programming concepts. This is an up-front cost but often leads to cleaner architecture (e.g. no shared-state locking issues to debug).

In conclusion, adopting Elixir/OTP for digital twins may involve new skills but offers a compelling trade-off: robust, scalable systems with fewer surprises under load.

Sources: Industry and academic references on digital twin definitions 【5†L41-L46】 , OTP documentation 【18†L18-L24】 【31†L18-L24】 【33†L18-L23】 【35†L14-L21】 【41†L12-L16】 , Erlang blog posts on IoT and OTP benefits 【15†L110-L119】 【15†L128-L137】 , and benchmark studies comparing Erlang/Go/Scala 【26†L19-L23】 have been used throughout this report to substantiate claims.
